


CPSC 221: Algorithms and Data Structures

Lecture #7

Sweet, Sweet Tree Hives



(B+-Trees, that is)

Steve Wolfman

2014W1

1

Today's Outline

- Addressing our other problem
- B+-tree properties
- Implementing B+-tree insertion and deletion
- Some final thoughts on B+-trees


2

M-ary Search Tree

- Maximum branching factor of M
- Complete tree has height $h \cong \log_M N$
- Each internal node in a complete tree has

$M - 1$ keys

runtime:




Actually, $h = \lceil \log_M (N(M - 1) + 1) \rceil - 1$

3

Incomplete M-ary Search Tree ☹


- Just like a binary tree, though, complete m-ary trees can store m^0 keys, $m^0 + m^1$ keys, $m^0 + m^1 + m^2$ keys, ...
- What about numbers in between??



4

B+-Trees

- B+-Trees are specialized M -ary search trees
- Each node has many keys
 - at least some minimum # of keys
 - subtree between two keys x and y contains values v such that $x \leq v < y$
 - binary search within a node to find correct subtree
- Each node takes one full {page, block, line} of memory
- ALL the leaves are at the same depth!



5

Today's Outline

- Addressing our other problem
- B+-tree properties
- Implementing B+-tree insertion and deletion
- Some final thoughts on B+-trees

6

B+-Tree Properties

- Properties
 - maximum branching factor of M
 - the root has between 2 and M children *or* at most L keys/values
 - other internal nodes have between $\lceil M/2 \rceil$ and M children
 - internal nodes contain only search keys (no data)
 - smallest datum between search keys x and y equals x
 - each (non-root) leaf contains between $\lceil L/2 \rceil$ and L keys/values
 - all leaves are at the same depth
- Result
 - tree is $\Theta(\log_M n)$ deep (between $\log_{M/2} n$ and $\log_M n$)
 - all operations run in $\Theta(\log_M n)$ time
 - operations get about $M/2$ to M or $L/2$ to L items at a time

B+-Tree Properties[‡]

- Properties
 - maximum branching factor of M
 - the root has between 2 and M children *or* at most L keys/values
 - other internal nodes have between $\lceil M/2 \rceil$ and M children
 - internal nodes contain only *search* keys (no data)
 - smallest datum between search keys x and y equals x
 - each (non-root) leaf contains between $\lceil L/2 \rceil$ and L keys/values
 - all leaves are at the same depth
- Result
 - tree is $\Theta(\log_M n)$ deep (between $\log_{M/2} n$ and $\log_M n$)
 - all operations run in $\Theta(\log_M n)$ time
 - operations get about $M/2$ to M or $L/2$ to L items at a time

[‡]These are B⁺-Trees. B-Trees store data at internal nodes.

B+-Tree Properties

- Properties
 - maximum branching factor of M
 - the root has between 2 and M children *or* at most L keys/values
 - other internal nodes have between $\lceil M/2 \rceil$ and M children
 - internal nodes contain only search keys (no data)
 - smallest datum between search keys x and y equals x
 - each (non-root) leaf contains between $\lceil L/2 \rceil$ and L keys/values
 - all leaves are at the same depth
- Result
 - tree is $\Theta(\log_M n)$ deep (between $\log_{M/2} n$ and $\log_M n$)
 - all operations run in $\Theta(\log_M n)$ time
 - operations get about $M/2$ to M or $L/2$ to L items at a time

B+-Tree Properties

- Properties
 - maximum branching factor of M
 - the root has between 2 and M children *or* at most L keys/values
 - other internal nodes have between $\lceil M/2 \rceil$ and M children
 - internal nodes contain only search keys (no data)
 - smallest datum between search keys x and y equals x
 - each (non-root) leaf contains between $\lceil L/2 \rceil$ and L keys/values
 - all leaves are at the same depth
- Result
 - height is $\Theta(\log_M n)$ between $\log_{M/2} (2n/L)$ and $\log_M (n/L)$
 - all operations run in $\Theta(\log_M n)$ time
 - operations get about $M/2$ to M or $L/2$ to L items at a time


Today's Outline

- Addressing our other problem
- B+-tree properties
- Implementing B+-tree insertion and deletion
- Some final thoughts on B+-trees


11

B+-Tree Nodes

- Internal node
 - i search keys; $i+1$ subtrees; $M - i - 1$ inactive entries




- Leaf
 - j data keys; $L - j$ inactive entries



12


Example

B+-Tree with $M = 4$
and $L = 4$



As with other dictionary data structures,
we show a version with no data, only keys, but *only* for simplicity!¹³

Making a B+-Tree




B-Tree with $M = 3$
and $L = 2$

Now, Insert(1)?

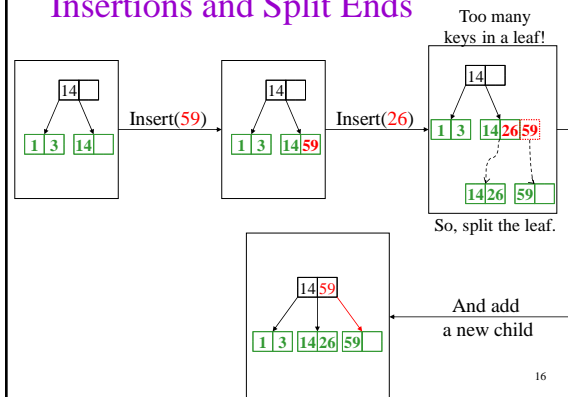
14

Splitting the Root



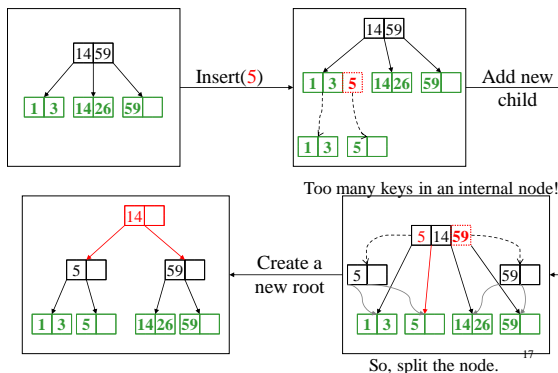
15

Insertions and Split Ends



16

Propagating Splits



17

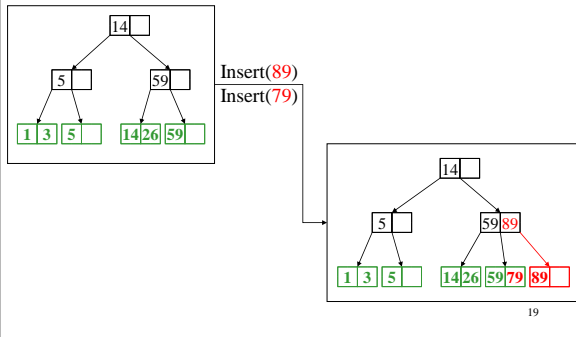
Insertion in Boring Text

- Insert the key in its leaf
- If the leaf ends up with $L+1$ items, **overflow!**
 - Split the leaf into two nodes:
 - original with $\lceil (L+1)/2 \rceil$ items
 - new one with $\lfloor (L+1)/2 \rfloor$ items
 - Add the new child to the parent
 - If the parent ends up with $M+1$ items, **overflow!**
- If an internal node ends up with $M+1$ items, **overflow!**
 - Split the node into two nodes:
 - original with $\lceil (M+1)/2 \rceil$ items
 - new one with $\lfloor (M+1)/2 \rfloor$ items
 - Add the new child to the parent
 - If the parent ends up with $M+1$ items, **overflow!**
- Split an overflowed root in two and hang the new nodes under a new root


This makes the tree deeper!

18

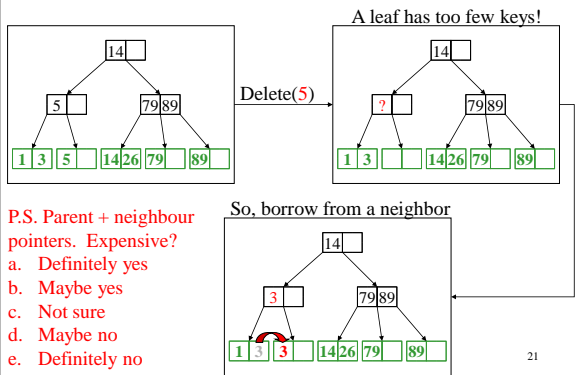
After More Routine Inserts




Deletion




Deletion and Adoption



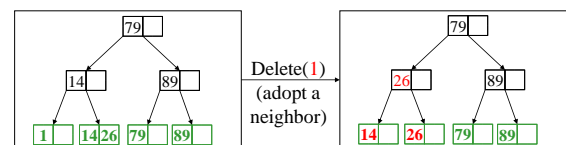
Deletion with Propagation



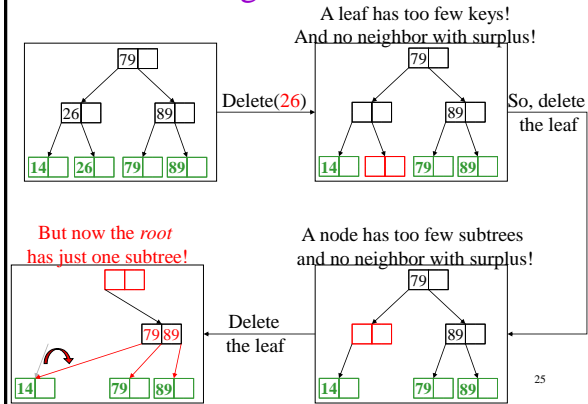
Finishing the Propagation (More Adoption)



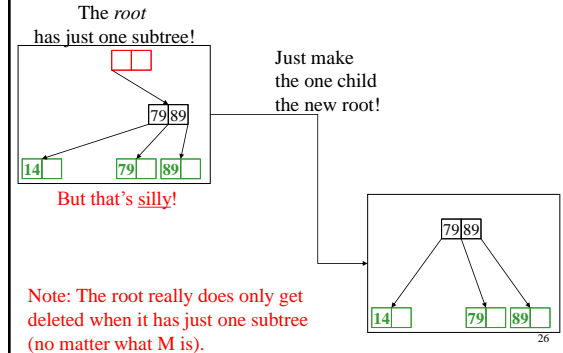
A Bit More Adoption



Pulling out the Root



Pulling out the Root (continued)



Deletion in Two Boring Slides of Text

- Remove the key from its leaf
- If the leaf ends up with fewer than $\lceil L/2 \rceil$ items, **underflow!**
 - Adopt data from a neighbor; update the parent
 - If borrowing won't work, delete node and divide keys between neighbors
 - If the parent ends up with fewer than $\lceil M/2 \rceil$ items, **underflow!**

Will dumping keys always work if adoption does not?

- Yes
- It depends
- No

27

Deletion Slide Two

- If a node ends up with fewer than $\lceil M/2 \rceil$ items, **underflow!**
 - Adopt subtrees from a neighbor; update the parent
 - If borrowing won't work, delete node and divide subtrees between neighbors
 - If the parent ends up with fewer than $\lceil M/2 \rceil$ items, **underflow!**
- If the root ends up with only one child, make the child the new root of the tree

This reduces the height of the tree!

28

Today's Outline

- Addressing our other problem
- B+-tree properties
- Implementing B+-tree insertion and deletion
- Some final thoughts on B+-trees

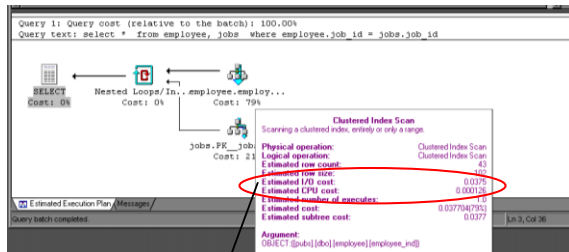
29

Thinking about B-Trees

- B+-Tree insertion can cause (expensive) splitting and propagation (could we do something like borrowing?)
- B+-Tree deletion can cause (cheap) borrowing or (expensive) deletion and propagation
- Propagation is rare if M and L are large (Why?)
- Repeated insertions and deletion can cause thrashing
- If $M = L = 128$, then a B-Tree of height 4 will store at least 30,000,000 items

30

Cost of a Database Query (14 years ago... more skewed now!)



I/O to CPU ratio is 300!

31

A Tree with Any Other Name

FYI:

- B+-Trees with $M = 3$, $L = x$ are called 2-3 trees
- B+-Trees with $M = 4$, $L = x$ are called 2-3-4 trees
- Other balanced trees include Red-Black trees (rotation-based), Splay Trees (rotation-based and *amortized* $O(\lg n)$ bounds), B-trees, B*-trees, ...

32

To Do

- Hashing readings

33

Coming Up

- In some order:
 - Everyone Gets a Crack at Parallelism
 - Hash Tables

34