

A Sophomoric Introduction to Shared-Memory Parallelism and Concurrency

Lecture 3 Parallel Prefix, Pack, and Sorting

Steve Wolfman, based on work by Dan Grossman

MOTIVATION: *Pack*

AKA, **filter**, like getting all elts less than the pivot in QuickSort.

Given an array **input**, produce an array **output** containing only elements such that **f(elt)** is **true**

Example: **input** [17, 4, 6, 8, 11, 5, 13, 19, 0, 24]
f: is elt > 10
output [17, 11, 13, 19, 24]

Parallelizable? Sure, using a list concatenation reduction.

Efficiently parallelizable on arrays?

Can we just put the output straight into the array *at the right spots*?

Sophomoric Parallelism and Concurrency, Lecture 3

2

MOTIVATION: *Pack as map, reduce, prefix combo??*

Given an array **input**, produce an array **output** containing only elements such that **f(elt)** is **true**

Example: **input** [17, 4, 6, 8, 11, 5, 13, 19, 0, 24]
f: is elt > 10

Which pieces can we do as maps, reduces, or prefixes?

Sophomoric Parallelism and Concurrency, Lecture 3

3

MOTIVATION: *Parallel prefix sum to the rescue (if only we had it!)*

1. Parallel map to compute a **bit-vector** for true elements
input [17, 4, 6, 8, 11, 5, 13, 19, 0, 24]
bits [1, 0, 0, 0, 1, 0, 1, 1, 0, 1]
2. Parallel prefix-sum on the bit-vector
bitsum [1, 1, 1, 1, 2, 2, 3, 4, 4, 5]
3. Parallel map to produce the output
output [17, 11, 13, 19, 24]

```
output = new array of size bitsum[n-1]
FORALL(i=0; i < input.size(); i++){
    if(bits[i])
        output[bitsum[i]-1] = input[i];
}
```

Sophomoric Parallelism and Concurrency, Lecture 3

4

The prefix-sum problem

Given a list of integers as input, produce a list of integers as output where **output[i] = input[0]+input[1]+...+input[i]**

Sequential version is straightforward:

```
vector<int> prefix_sum(const vector<int>& input){
    vector<int> output(input.size());
    output[0] = input[0];
    for(int i=1; i < input.size(); i++){
        output[i] = output[i-1]+input[i];
    }
    return output;
}
```

Example:

input	42	3	4	7	1	10
output						

Sophomoric Parallelism and Concurrency, Lecture 3

5

The prefix-sum problem

Given a list of integers as input, produce a list of integers as output where **output[i] = input[0]+input[1]+...+input[i]**

Sequential version is straightforward:

```
Vector<int> prefix_sum(const vector<int>& input){
    vector<int> output(input.size());
    output[0] = input[0];
    for(int i=1; i < input.size(); i++){
        output[i] = output[i-1]+input[i];
    }
    return output;
}
```

Why isn't this (obviously) parallelizable? Isn't it just map or reduce?

Work:

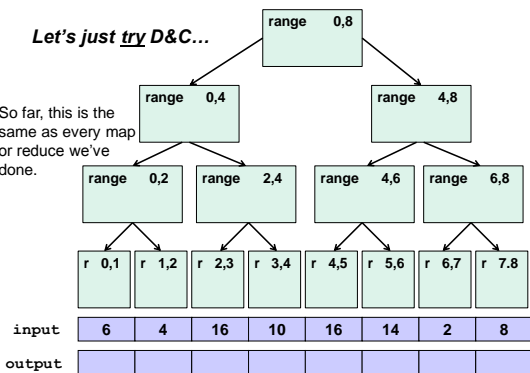
Span:

Sophomoric Parallelism and Concurrency, Lecture 3

6

Let's just try D&C...

So far, this is the same as every map or reduce we've done.

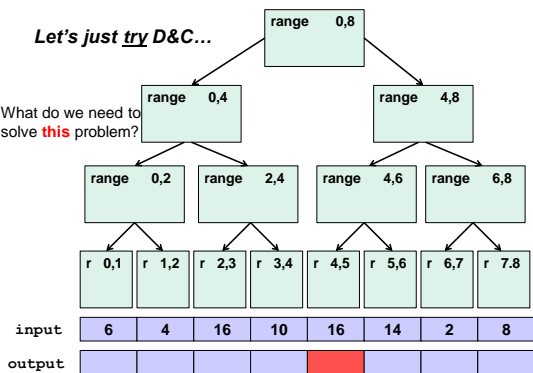


Sophomore Parallelism and Concurrency, Lecture 3

7

Let's just try D&C...

What do we need to solve **this** problem?

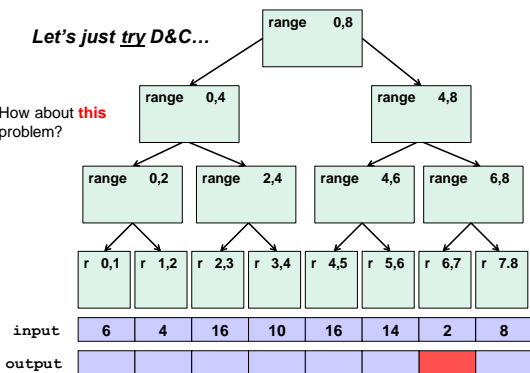


Sophomore Parallelism and Concurrency, Lecture 3

8

Let's just try D&C...

How about **this** problem?



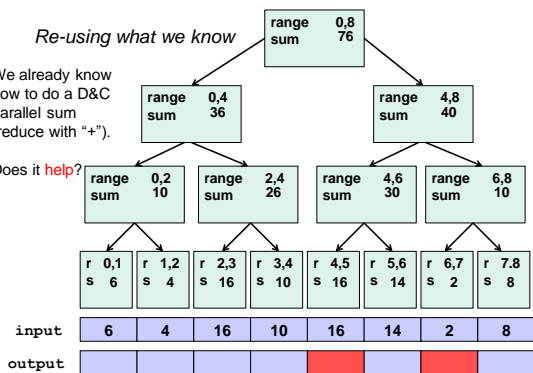
Sophomore Parallelism and Concurrency, Lecture 3

9

Re-using what we know

We already know how to do a D&C parallel sum (reduce with "+").

Does it help?

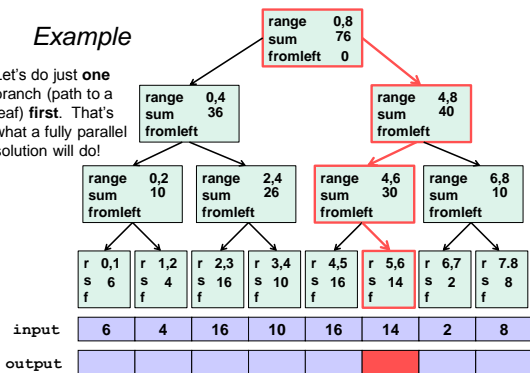


Sophomore Parallelism and Concurrency, Lecture 3

10

Example

Let's do just one branch (path to a leaf) first. That's what a fully parallel solution will do!



Sophomore Parallelism and Concurrency, Lecture 3

Algorithm from [Ladner and Fischer, 1977]

Parallel prefix-sum

The parallel-prefix algorithm does two passes:

1. build a "sum" tree bottom-up
2. traverse the tree top-down, accumulating the sum from the left

Sophomore Parallelism and Concurrency, Lecture 3

12

The algorithm, step 1

- Step one does a parallel sum to build a binary tree:
 - Root has sum of the range $[0, n]$
 - An internal node with the sum of $[lo, hi)$ has
 - Left child with sum of $[lo, middle)$
 - Right child with sum of $[middle, hi)$
 - A leaf has sum of $[i, i+1)$, i.e., `input[i]` (or an appropriate larger region w/a cutoff)

How? Parallel sum but explicitly build a tree:

```
return left+right;  =>  return new Node(left->sum + right->sum,
                                left, right);
```

Step 1: Work? Span?

Sophomoric Parallelism and Concurrency, Lecture 3

13

The algorithm, step 2

- Parallel map, passing down a `fromLeft` parameter
 - Root gets a `fromLeft` of 0
 - Internal nodes pass along:
 - to its left child the same `fromLeft` (already calculated in step 1!)
 - to its right child `fromLeft` plus its left child's sum
 - At a leaf node for array position i ,
`output[i]=fromLeft+input[i]`

How? A map down the step 1 tree, leaving results in the output array.

Notice the invariant: `fromLeft` is the sum of elements left of the node's range

Step 2: Work? Span?

Sophomoric Parallelism and Concurrency, Lecture 3

14

Parallel prefix-sum

The parallel-prefix algorithm does two passes:

- build a "sum" tree bottom-up
- traverse the tree top-down, accumulating the sum from the left

Step 1: Work: $O(n)$ Span: $O(\lg n)$
 Step 2: Work: $O(n)$ Span: $O(\lg n)$

Overall: Work? Span?

Parallelism (work/span)?

Sophomoric Parallelism and Concurrency, Lecture 3

In practice, of course, we'd use a sequential cutoff! 15

Parallelizing Quicksort

Recall quicksort was sequential, in-place, expected time $O(n \lg n)$

- | | |
|--|----------------------------------|
| | Best / expected case work |
| 1. Pick a pivot element | $O(1)$ |
| 2. Partition all the data into: | $O(n)$ |
| A. The elements less than the pivot | |
| B. The pivot | |
| C. The elements greater than the pivot | |
| 3. Recursively sort A and C | $2T(n/2)$ |

How do we parallelize this?

What span do we get?

$T_\infty(n) =$

Sophomoric Parallelism and Concurrency, Lecture 3

16

Parallelizing Quicksort

Recall quicksort was sequential, in-place, expected time $O(n \lg n)$

- | | |
|--|----------------------------------|
| | Best / expected case work |
| 1. Pick a pivot element | $O(1)$ |
| 2. Partition all the data into: | $O(n)$ |
| A. The elements less than the pivot | |
| B. The pivot | |
| C. The elements greater than the pivot | |
| 3. Recursively sort A and C | $2T(n/2)$ |

How should we parallelize this?

Parallelize the recursive calls as we usually do in fork/join D&C.

Parallelize the partition by doing two packs (filters) instead.

Sophomoric Parallelism and Concurrency, Lecture 3

17

Analyzing $T_\infty(n) = \lg n + T_\infty(n/2)$

Turns out our techniques from way back at the start of the term will work just fine for this:

$$T_\infty(n) = \lg n + T_\infty(n/2) \quad \text{if } n > 1$$

$$= 1 \quad \text{otherwise}$$

Sophomoric Parallelism and Concurrency, Lecture 3

18